

# Bias, Censorship, and Accuracy: AI Models for Online Moderation

Josef Karpinski (jjk21004)  
Alexander Manos (agm21019)  
Samuel Cultrera (sac19019)  
Malcolm Ferguson (mrf20004)




# Project Overview

- Compared two AI models in their ability to classify online text comments as toxic or non-toxic
- Trained models using the Jigsaw Toxic Comment Classification Dataset

We discussed...

- Ethical concerns related to automatic content moderation and free speech
- Practical use and real world implications of these algorithms






# Exploring the Dataset



# Jigsaw Toxic Comment Classification Dataset



## Purpose

- Widely used to develop machine learning models that aim to classify comments as toxic or non-toxic

## Why We Chose This Dataset


- Well-labeled data, ideal for supervised learning
- Real world use case



# Specifications

- Number of Samples: ~160,000
- Data-Type: English text comments from wikipedia talk pages
- Format: Multi-label classification
  - 6 Labels - toxic, severe toxic, obscene, threat, insult, identity hate

Label	Positive	Negative	Proportion
toxic	15294	144277	10.6%
severe toxic	1595	157976	1.0%
obscene	8449	151122	5.6%
threat	478	159093	0.3%
insult	7877	151694	5.2%
identity hate	1405	158166	0.9%
total	15294	144277	10.6%



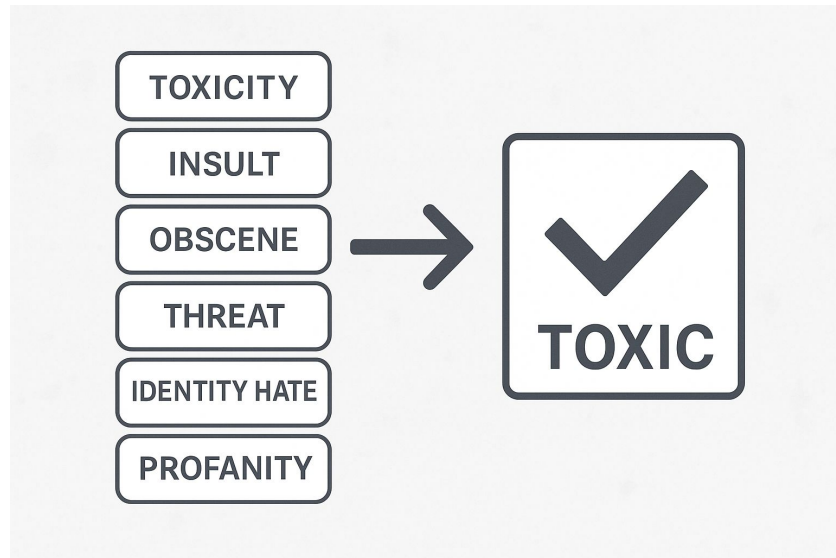


# Methodology



# Combining Labels


- For simplicity, we will combine our 6 labels in one binary label
- We will consider a text sample toxic if just 1 of the 6 labels is positive
- This turns the problem into binary classification, instead of multi classification






# Logistic Regression

- Splitting the dataset
  - We split the dataset using an 80-20 split, meaning 80% training and 20% testing
- Vectorization of text data
  - Since we are dealing with text data, it is important to vectorize the data before we input it into our model
  - Use TfidfVectorizer from sklearn.feature\_extraction.text
- Defining the model and prediction
  - Use LogisticRegression from sklearn.linear\_model



## LOGISTIC REGRESSION







# Logistic Regression (Code)

- Splitting the dataset

```
X_train, X_test, y_train, y_test = train_test_split(  
    df['comment_text'], # input  
    df['toxic_any'], # output  
    test_size=0.2, # 20% test, 80% train  
    random_state=42  
)
```

- Vectorization of text data


```
vectorizer = TfidfVectorizer(stop_words='english', max_features=5000)  
X_train_vec = vectorizer.fit_transform(X_train)  
X_test_vec = vectorizer.transform(X_test)
```

- Defining the model and prediction

```
log_reg_model = LogisticRegression(max_iter=1000)  
log_reg_model.fit(X_train_vec, y_train)  
y_pred_log_reg = log_reg_model.predict(X_test_vec)
```



# Large Language Model



- Unbiased Toxic RoBERTa Multiclassifier
  - Hugging Face: <https://huggingface.co/unitary/unbiased-toxic-roberta>
  - Trained on our dataset and a few others
- Sampling the dataset
  - Classification with an LLM is quite slow, so we drew a subset of 100 samples to do testing
- Testing and Classification
  - Run the LLM on each sample to get score for all of the LLM labels
  - If any of our desired scores are above a threshold, we classify the text as toxic
  - We must define our own classification function



# LLM (Code 1)

- Downloading the classifier

```
llm_classifier = pipeline("text-classification", model="unitary/unbiased-toxic-roberta")
```

- Dataset Sampling

```
df_llm = df.sample(1000, random_state=42)

X_llm = df_llm['comment_text']
y_llm = df_llm['toxic_any']
```

## LLM Labels:

- toxicity
- severe\_toxicity
- obscene
- identity\_attack
- insult
- threat
- sexual\_explicit



## LLM (Code 2)

- Custom function for classification

```
def classify_llm(text, threshold=0.5):  
    try:  
        predictions = llm_classifier(text, top_k=None)  
        # see if any of our labels are above the threshold  
        for pred in predictions:  
            if pred['label'] in TOXIC_LABELS:  
                if pred['score'] > threshold:  
                    return 1  
        # if here, none of the labels were sufficient  
        return 0  
    except:  
        return 0
```

```
TOXIC_LABELS = {  
    "toxicity",  
    "severe_toxicity",  
    "obscene",  
    "threat",  
    "insult",  
    "identity_attack",  
    "sexual_explicit"  
}
```

- Run LLM on test data

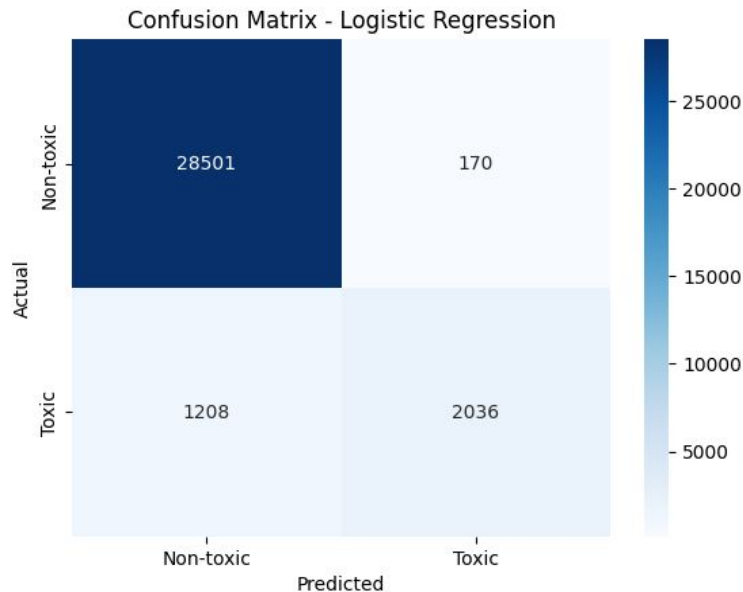
```
y_pred_llm = X_llm.apply(classify_llm, threshold=0.5)
```



# Results and Analysis

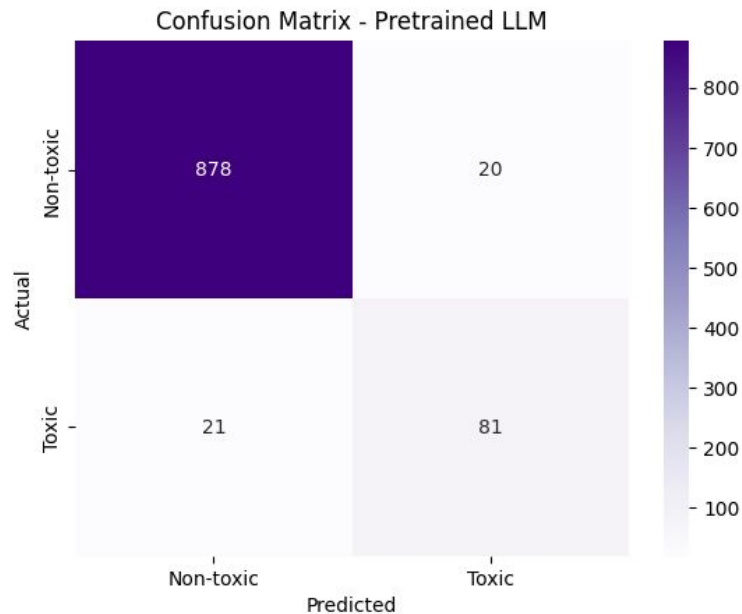
# Results - Our Model

- Results from running on test data:
  - Precision: 0.9229
  - Recall: 0.6276
  - F1-score: 0.7473
- Overall accuracy = 0.9568




# Results - LLM

- Results from running on test data:
  - Precision: 0.8020
  - Recall: 0.7941
  - F1-score: 0.7473
- Overall accuracy = 0.9590



# Results - Comparison

- Overall accuracy very similar
- LLM has a 17% higher recall
  - More false positives, but less false negatives
  - How should this be balanced?
- Must take data balance into account when analyzing accuracy







# Ethical Discussion






# Accuracy

It is Very Difficult for Artificial Intelligence to Fully Interpret Messages

AI often lacks nuance  
and cannot process  
context such as sarcasm  
Risk of false negatives  
or false positives





# Privacy Vs Protection



Since AI Cannot Reach 100% Accuracy, a Choice Must Be Made

## Under Flagging

- ❖ Preserves freedom of expression
- ❖ Reduces false positives
- ❖ Allows harmful content to persist
- ❖ Erodes platform trust


## Over Flagging

- ❖ Protects users
- ❖ May restrict free speech
- ❖ Suppresses legitimate expression
- ❖ Frustrates users




# Potential Solution

AI Algorithm



Human Review



Official Decision





**Questions?**



# References

Jigsaw Toxic Comment Classification Dataset:

<https://www.kaggle.com/competitions/jigsaw-toxic-comment-classification-challenge/data>

Unbiased Toxic RoBERTa Pre-trained LLM:

<https://huggingface.co/unitary/unbiased-toxic-roberta>

GitHub Repository

<https://github.com/josef-karpinski/content-moderation-cse3000>